

---

# Solving OpenAI’s Car Racing Environment with Deep Reinforcement Learning and Dropout

---

**Patrik Gerber\***

University of Oxford  
patrik.gerber@ccc.ox.ac.uk

**Jiajing Guan\***

George Mason University  
jiajingguan@gmail.com

**Elvis Nunez\***

Brown University  
elvis@brown.io

**Kaman Phamdo\***

University of Maryland  
kaman@phamdo.com

**Nicholas Malaya**

AMD Research  
nicholas.malaya@amd.com

**Tonmoy Monsoor**

UCLA  
mtonmoy@g.ucla.edu

## 1 Introduction

Deep reinforcement learning methods have proven successful in solving well-defined computer games. In particular, the DDQN algorithm has achieved superhuman performance in several Atari games [4]. However, deep RL agents are limited by their sensitivity to perturbations in the environment. They can fail catastrophically when applied to environments that differ even slightly from where they were trained.

The OpenAI Gym’s CarRacing-v0 environment [3] is a top-down view car racing game where the goal is to visit all tiles of the racetrack as quickly as possible. The agent receives a reward of  $-0.1$  at each time step and  $\frac{1000}{N}$  for each track tile visited, where  $N$  is the total number of tiles in the racetrack. The state is represented by  $96 \times 96$  RGB screenshots of the game. The game ends when the agent visits all tiles on the track or when 1000 frames have passed. This environment is considered solved by OpenAI’s guidelines when the agent achieves an average score of 900 or above over 100 consecutive games. This game is particularly challenging because the racetrack is randomly generated, which introduces stochasticity and complexity into the environment. The challenge has been solved using a generative neural network model [1], but previous attempts using deep RL methods have not been successful [2].

In this paper, we describe a deep RL algorithm that uses a convolutional architecture with dropout that successfully solves the game. Remarkably, the model is trained on a limited environment made up of 3 tracks. This result shows that regularization methods such as dropout can mitigate the overfit usually exhibited by deep RL. To the best of our knowledge, this solution is the first to solve the challenge using deep RL.

## 2 Method

We adapted the DDQN algorithm [5] for our approach. The input to the Q-network is a  $96 \times 96 \times 4$  image produced by stacking 4 consecutive frames of the game while the output is five dimensional, corresponding to the elements of the discretized action space. As proposed in the original DQN algorithm [4], the architecture for the Q-network consists of 3 convolutional layers followed by 2 dense layers and an output for each of the 5 actions. Each hidden layer is followed by a rectified nonlinearity. Dropout is applied to the second convolutional layer. We train over a maximum of 3000 episodes and use early stopping to ensure that the best performing model is selected for analysis.

We trained several models with varying observability of the state space. First, we train the model on a single racetrack. Second, we randomly alternate between 3 fixed racetracks. Third, the agent is shown random racetracks from the state space. All models are tested on random racetracks. We

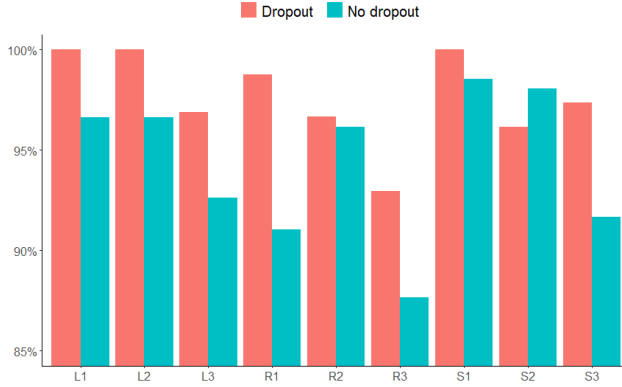


Figure 1: Proportion of tiles visited on different types of curves.




Figure 2: Different types of curves.

performed a baseline experiment with no dropout and experimented with the dropout probability. The results from the best performing models are displayed in Table 1.

### 3 Performance Analysis

As shown in Table 1, the average scores over 100 random games improved in all three environments when dropout was applied to the Q-network. This suggests that dropout acts as an effective regularizer by improving the agent’s performance in situations it has not observed during training. In order to investigate the generalizability of these models, we developed a curve classification algorithm. Each curve in the racetrack is characterized as either a Left, Right, or S-shaped curve. Then, the steepness of the curve is ranked from 1 to 3, where 1 represents a shallow curve and 3 represents a steep curve. For each of these different types of curves, we record the percentage of tiles cleared.

Environment	Dropout	Average Score
1 Fixed Track	None	$849.99 \pm 78.72$
3 Fixed Tracks	None	$853.88 \pm 127.71$
Random Tracks	None	$854.83 \pm 107.14$
1 Fixed Track	0.7	$894.38 \pm 24.5$
3 Fixed Tracks	0.5	$906.67 \pm 23.6$
Random Tracks	0.5	$892.62 \pm 41.48$

Table 1: Average scores over 100 consecutive games on random racetracks. Test tracks were generated from random seeds 1 through 100.

Figure 1 compares the performance of the models trained on a single racetrack. Even though both models were trained on the same racetrack, the model with dropout performed better than the model without dropout, especially on curves that the agent did not see examples of during training. This indicates that using dropout is an effective regularizer and improves the generalizability of these models.

### 4 Conclusion

We demonstrated a deep RL solution to CarRacing-v0 where the agent was trained on a limited subset of the state space. Even with partial knowledge about the environment, the agent is capable of making informed decisions. By applying dropout as a regularization technique, we improved the algorithm’s ability to generalize to situations it has not observed during training. These results suggest that widespread regularization techniques such as dropout could be successfully applied to deep RL algorithms to make them more robust under partial observability of the environment.

## References

- [1] David Ha and Jürgen Schmidhuber. World models. *CoRR*, abs/1803.10122, 2018.
- [2] M. A. Farhan Khan and Oguz H. Elibol. Car racing using reinforcement learning. 2016.
- [3] Oleg Klimov. CarRacing-v0. <https://gym.openai.com/envs/CarRacing-v0/>. Accessed: 2018-09-01.
- [4] Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Andrei A Rusu, Joel Veness, Marc G Bellemare, Alex Graves, Martin Riedmiller, Andreas K Fidjeland, Georg Ostrovski, et al. Human-level control through deep reinforcement learning. *Nature*, 518(7540):529, 2015.
- [5] Hado Van Hasselt, Arthur Guez, and David Silver. Deep reinforcement learning with double q-learning. In *AAAI*, volume 2, page 5. Phoenix, AZ, 2016.

## 5 Appendix

### 5.1 Network architecture

See the table below for the exact architecture used in our models.

Layer	Topology	Activation
1	Convolutional, 32 8x8 kernels with stride 4	ReLU
2	Convolutional, 64 4x4 kernels with stride 2	ReLU
3	Convolutional, 64 3x3 kernels with stride 1	ReLU
4	Dense, 512 neurons	ReLU
output	Dense, 5 neurons	Linear

### 5.2 Action space

The action space of OpenAI’s CarRacing environment is the set  $[-1, 1] \times [0, 1]^2 \subset \mathbb{R}^3$ . We discretize this space into five possible actions. See the table below for more detail.

Value	Interpretation
$[-1, 0, 0]$	Steer left
$[1, 0, 0]$	Steer right
$[0, 1, 0]$	Accelerate
$[0, 0, 0.8]$	Decelerate
$[0, 0, 0]$	Do nothing

### 5.3 Training parameters

See the table below for the training parameters that were used during training for our models.

Parameter	Value
Exploration rate	Decrease from 1 to 0.1 linearly over the first 250000 frames, 0.1 thereafter
Optimizer	Default Tensorflow AdamOptimizer
Learning rate	0.00025
Discount value	0.99
Batch size	32
Replay capacity	100000